

In [6]:

```
from termcolor import colored
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import os

from sklearn.metrics import roc_auc_score
from sklearn.metrics import roc_curve, auc
import matplotlib.pyplot as plt
import IPython
%matplotlib inline

import keras
from keras.layers import Input, Dense, Dropout
from keras.layers import BatchNormalization
from keras.models import Model
from keras.optimizers import Adam
from keras import backend
from keras.models import load_model
from keras.layers import Dropout, Masking
from keras.preprocessing.text import Tokenizer
from keras.preprocessing.text import hashing_trick
from keras.preprocessing.text import text_to_word_sequence
from sklearn.model_selection import train_test_split
```

The following presentation will demonstrate the use of two ML approaches for solving a classification problem.

The first model is a simple fully connected logistic regression model.

The second model is one I'm particularly excited about, the Generative Adversarial Network (GAN).

GAN's are constructed by pitting two networks against each other in a zero sum game. The Discriminator is trained to detect fake loan applications and labels whether an application is likely to be approved or rejected. The Generator is then fed noise to produce fake loan applications to the Discriminator with the objective of passing as a valid loan. As the Discriminator learns to detect fake applications it forces the Generator to learn more clever ways to pass a fake. Essentially, the Generator will learn the fundamental distribution of the dataset. Over time, each network reinforces the others weights to achieve an optimal equilibrium.

It is worth nothing that GAN's have been successful in several areas including datasets with limited labelled data and image generation.

Let's start by downloading and loading all HMDA data filtered by state of Connecticut for 2014-2016 limited to those intended for home purchase.

In [221]:

```
# Note: link to data and filter settings
# Note: https://www.consumerfinance.gov/data-research/hmda/explore#!/as_of_year=201
data = pd.read_csv('/home/shant/Downloads/hmda/hmda_lar.csv', low_memory=False, header=0)
data.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 171251 entries, 0 to 171250
Data columns (total 78 columns):
action_taken                171251 non-null int64
action_taken_name           171251 non-null object
agency_code                 171251 non-null int64
agency_abbr                 171251 non-null object
agency_name                 171251 non-null object
applicant_ethnicity         171251 non-null int64
applicant_ethnicity_name    171251 non-null object
applicant_income_000s       148174 non-null float64
applicant_race_1            171251 non-null int64
applicant_race_2            648 non-null float64
applicant_race_3            36 non-null float64
applicant_race_4            8 non-null float64
applicant_race_5            9 non-null float64
applicant_race_name_1       171251 non-null object
applicant_race_name_2       648 non-null object
applicant_race_name_3       36 non-null object
applicant_race_name_4       8 non-null object
applicant_race_name_5       9 non-null object
applicant_sex               171251 non-null int64
applicant_sex_name          171251 non-null object
application_date_indicator  171251 non-null int64
as_of_year                  171251 non-null int64
census_tract_number        170932 non-null float64
co_applicant_ethnicity      171251 non-null int64
co_applicant_ethnicity_name 171251 non-null object
co_applicant_race_1         171251 non-null int64
co_applicant_race_2         192 non-null float64
co_applicant_race_3         8 non-null float64
co_applicant_race_4         2 non-null float64
co_applicant_race_5         1 non-null float64
co_applicant_race_name_1    171251 non-null object
co_applicant_race_name_2    192 non-null object
co_applicant_race_name_3    8 non-null object
co_applicant_race_name_4    2 non-null object
co_applicant_race_name_5    1 non-null object
co_applicant_sex            171251 non-null int64
co_applicant_sex_name       171251 non-null object
county_code                 171051 non-null float64
county_name                 171051 non-null object
denial_reason_1             11674 non-null float64
denial_reason_2             2545 non-null float64
denial_reason_3             455 non-null float64
denial_reason_name_1        11674 non-null object
denial_reason_name_2        2545 non-null object
denial_reason_name_3        455 non-null object
edit_status                 21002 non-null float64
edit_status_name            21002 non-null object
hoepa_status                171251 non-null int64
hoepa_status_name           171251 non-null object
```

```

lien_status          171251 non-null int64
lien_status_name     171251 non-null object
loan_purpose           171251 non-null int64
loan_purpose_name      171251 non-null object
loan_type            171251 non-null int64
loan_type_name       171251 non-null object
msamd                162077 non-null float64
msamd_name           162077 non-null object
owner_occupancy      171251 non-null int64
owner_occupancy_name 171251 non-null object
preapproval          171251 non-null int64
preapproval_name     171251 non-null object
property_type        171251 non-null int64
property_type_name   171251 non-null object
purchaser_type       171251 non-null int64
purchaser_type_name  171251 non-null object
respondent_id        171251 non-null object
sequence_number      171251 non-null int64
state_code           171251 non-null int64
state_abbr           171251 non-null object
state_name           171251 non-null object
hud_median_family_income 170932 non-null float64
loan_amount_000s     171251 non-null int64
number_of_1_to_4_family_units 170932 non-null float64
number_of_owner_occupied_units 170903 non-null float64
minority_population   170932 non-null float64
population            170932 non-null float64
rate_spread          4653 non-null float64
tract_to_msamd_income 170927 non-null float64
dtypes: float64(23), int64(21), object(34)
memory usage: 101.9+ MB

```

At first glance we notice that most of the data has a scalar value and an associated label denoted by column names ending with "_name".

There are several fields indicating the reason for denying the application. We will exclude these features since that information is relevant after the target value has been determined.

Respondent ID stands out as valuable information related to the mortgage issuer. Let's convert these text columns to scalar values by using pandas factorize method.

Let's remove all the redundant name data and define our features.

Let's also look at the dataframe description to see if we should pre-normalize the data or use Batch Normalization instead.

In [224]:

```
data['issuer'] = data[['respondent_id']].apply(lambda col: pd.factorize(col)[0])

features = [
    'agency_code', 'applicant_ethnicity', 'applicant_income_000s', 'applicant_race_1',
    'applicant_race_4', 'applicant_race_5', 'applicant_sex', 'application_date_indicat
    'co_applicant_ethnicity', 'co_applicant_race_1', 'co_applicant_race_2', 'co_applic
    'co_applicant_race_5', 'co_applicant_sex', 'county_code',
    'edit_status',
    'hoepa_status', 'hud_median_family_income', 'lien_status', 'loan_amount_000s', 'loa
    'msamd', 'number_of_1_to_4_family_units', 'number_of_owner_occupied_units', 'owner
    'property_type', 'purchaser_type', 'rate_spread',
    'issuer',
    'sequence_number', 'state_code', 'tract_to_msamd_income']
```

Next, let's explore the different values associated with `actions_taken` in the distribution and define our target Y. There are 8 unique responses in the dataset. Since we are only interested in predicting which applications will be approved let's consolidate the 8 values into 0 for denial, 1 for approval.

For the purposes of this exercise we will assume that the following actions will be considered an approval:

1: Loan originated

We will assume the following actions will be considered denial:

3: Application denied by financial institution

We will drop all other actions since it is not explicit with respect to the objective function.

In [225]:

```
for k,v in data.groupby(['action_taken', 'action_taken_name']):
    print(k, "\t Count:", len(v))
data = data[data.action_taken.isin([1,3])]
data['approved'] = data.action_taken.isin([1]).astype(int)
data = data.fillna(0)

print("")
print("Y values:")
for k,v in data.groupby(['approved', 'action_taken_name']):
    print(k, "\t Count:", len(v))
```

```
(1, 'Loan originated') :          Count: 99210
(3, 'Application denied by financial institution') :      Count: 14220
```

Y values:

```
(0, 'Application denied by financial institution') :      Count: 14220
(1, 'Loan originated') :          Count: 99210
```

Let's explore the feature description

In [226]:

```
pd.set_option('display.float_format', lambda x: '%.3f' % x)
# print("---- Min ----")
# print(data[features].min())
# print("---- Max ----")
# print(data[features].max())
print(data[features].describe())
```

	agency_code	applicant_ethnicity	applicant_income_000s	\
count	113430.000	113430.000	113430.000	
mean	6.559	2.021	129.000	
std	2.297	0.462	213.266	
min	1.000	1.000	0.000	
25%	5.000	2.000	56.000	
50%	7.000	2.000	85.000	
75%	9.000	2.000	139.000	
max	9.000	4.000	9999.000	

	applicant_race_1	applicant_race_2	applicant_race_3	applicant_race_4	\
count	113430.000	113430.000	113430.000	113430.000	113
mean	4.825	0.019	0.001	0.000	
std	0.910	0.297	0.071	0.030	
min	1.000	0.000	0.000	0.000	
25%	5.000	0.000	0.000	0.000	
50%	5.000	0.000	0.000	0.000	
75%	5.000	0.000	0.000	0.000	
max	7.000	5.000	5.000	5.000	

	applicant_race_5	applicant_sex	application_date_indicator	\
count	113430.000	113430.000	113430.000	
mean	0.000	1.445	0.000	
std	0.035	0.646	0.000	
min	0.000	1.000	0.000	
25%	0.000	1.000	0.000	
50%	0.000	1.000	0.000	
75%	0.000	2.000	0.000	
max	5.000	4.000	0.000	

	...	owner_occupancy	population	preapproval
count	...	113430.000	113430.000	113430.000
mean	...	1.082	4932.893	2.501
std	...	0.291	1580.722	0.599
min	...	1.000	0.000	1.000
25%	...	1.000	3741.000	2.000
50%	...	1.000	4810.000	3.000

75%	...	1.000	5974.000	3.000
max	...	3.000	10289.000	3.000

	property_type	purchaser_type	rate_spread	issuer	sequence
_number \					
count	113430.000	113430.000	113430.000	113430.000	113
430.000					
mean	1.015	3.127	0.082	86.700	55
185.072					
std	0.160	3.217	0.433	112.092	137
275.915					
min	1.000	0.000	0.000	0.000	
1.000					
25%	1.000	0.000	0.000	16.000	
747.000					
50%	1.000	2.000	0.000	51.000	4
534.000					
75%	1.000	6.000	0.000	114.000	28
354.000					
max	3.000	9.000	8.530	671.000	1241
460.000					

	state_code	tract_to_msamd_income
count	113430.000	113430.000
mean	9.000	114.261
std	0.000	40.851
min	9.000	0.000
25%	9.000	89.540
50%	9.000	109.730
75%	9.000	134.080
max	9.000	256.740

[8 rows x 41 columns]

The following code blocks are intended to be executed during/after training to evaluate performance. We will load the saved test results from disk and compute basic statistics.

Let's analyze the results!

In [154]:

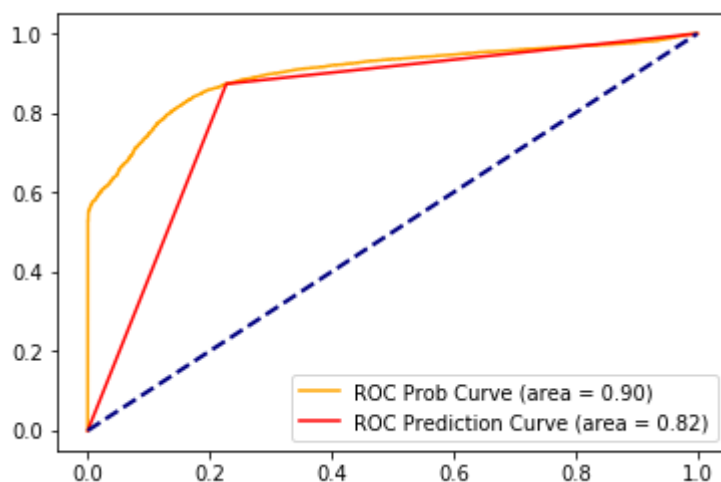
```
def analyze(filename):
    results = pd.read_csv('/home/shant/PycharmProjects/OWL/'+filename, low_memory=False)
    mislabeled = results[results.approved != results.prediction]
    total = len(results)
    error = len(mislabeled)/total
    accuracy = 1 - error
    denied = results.loc[results.approved==0, 'approved'].count()
    approved = results.loc[results.approved==1, 'approved'].count()
    denied_tp = results.loc[(results.approved==0)&(results.prediction==0), 'prediction'].count()
    approved_tp = results.loc[(results.approved==1)&(results.prediction==1), 'prediction'].count()
    roc_prob = roc_auc_score(results.approved, results.probs)
    roc_pred = roc_auc_score(results.approved, results.prediction)
    fpr, tpr, thresh = roc_curve(results.approved, results.probs)
    fpr_pred, tpr_pred, thresh_pred = roc_curve(results.approved, results.prediction)
    roc_auc = auc(fpr, tpr)
    roc_auc_pred = auc(fpr_pred, tpr_pred)
    print("Errors: \t {:.2f}%".format(error*100))
    print("Accuracy: \t {:.2f}%".format(accuracy*100))
    print("Denied \t", denied)
    print("Approved \t", approved)
    print("Denied True Pos \t {} / {:.2f}%".format(denied_tp, 100*(denied_tp/denied)))
    print("Approved True Pos \t {} / {:.2f}%".format(approved_tp, 100*(approved_tp/approved)))
    print("AUC Probs. \t {:.2f}%".format(roc_prob*100))
    print("AUC Predictions \t {:.2f}%".format(roc_pred*100))
    plt.plot(fpr, tpr, color='orange', label='ROC Prob Curve (area = %0.2f)' % roc_auc)
    plt.plot(fpr_pred, tpr_pred, color='red', label='ROC Prediction Curve (area = %0.2f)' % roc_auc_pred)
    plt.plot([0, 1], [0, 1], color='navy', lw=2, linestyle='--')
    plt.legend()
```

First, the simple logistic regression model

In [227]:

```
analyze('results_best.txt')
```

Errors:	13.89%
Accuracy:	86.11%
Denied	2788
Approved	19898
Denied True Pos	2154 / 77.26%
Approved True Pos	17381 / 87.35%
AUC Probs.	90.01
AUC Predictions	82.31

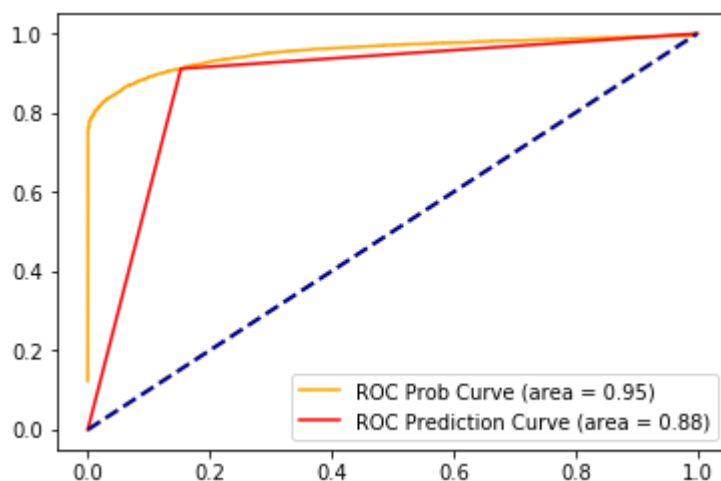


Let's do the same analysis for the GAN

In [233]:

```
analyze('results_gan_best.txt')
```

Errors:	9.64%
Accuracy:	90.36%
Denied	2788
Approved	19898
Denied True Pos	2362 / 84.72%
Approved True Pos	18136 / 91.14%
AUC Probs.	95.45
AUC Predictions	87.93



The GAN stands out in this exercise for several reason.

AUC and accuracy are far more stable during training than the logistic regression

While GAN's can be difficult/time consuming to train, it surpasses the logistic regression from the start